

ECS 174: Intro to Computer Vision, Spring 2020

Problem Set 4

Instructor: Yong Jae Lee (yongjaelee@ucdavis.edu)
Instructor: Krishna Kumar Singh (krsingh@ucdavis.edu)
TA: Haotian Liu (lhtliu@ucdavis.edu)
TA: Utkarsh Ojha (uojha@ucdavis.edu)
TA: Yuheng Li (yhnli@ucdavis.edu)

Due: Wednesday, June 10th, 11:59 PM

Instructions

1. Answer sheets must be submitted on Canvas. Hard copies will not be accepted.
2. Please submit your answer sheet containing the written answers in a file named: `FirstName_LastName_PS4.pdf`.
3. **You may complete the assignment individually or with a partner (i.e., maximum group of 2 people). If you worked with a partner, provide the name of your partner. We will be checking for instances of plagiarism/cheating.**

1 Short answer problems [60 points]

1. In the boosting algorithm AdaBoost, an ensemble of classifiers are selected in sequence, and the weight w_i on each labeled training example x_i is adjusted from iteration to iteration. What is the purpose of these weights? How do they influence which feature+classifier combination is selected in the next round? [10 points]
2. Invariant interest point detection and geometric verification using a parametric transformation are both commonly used in “instance” recognition (e.g., to recognize a landmark building in a tourist photo), but not in object category recognition (e.g., to recognize any dog). Explain why. [10 points]
3. How does a k -nearest neighbor classifier use the k nearest neighbors to make a label prediction? [10 points]
4. A deep neural network has multiple layers with non-linear activation functions (e.g., ReLU) in between each layer, which allows it to learn a complex non-linear function. Suppose instead we had a deep neural network without any non-linear activation functions. Concisely describe what effect this would have on the network. (Hint: can it still be considered a deep network?) [10 points]

5. One module of a standard convolutional neural network is the max-pooling operation. Given the 4x4 image below, perform max-pooling with a stride of 2 and a pooling window of 2x2. [10 points]

2	5	3	56
20	5	1	32
3	3	7	46
4	3	12	23

6. Determine whether each statement is true or false (no need to explain your answer) [10 points]:
- To detect profile (side) views of faces with the boosting-based Viola-Jones face detector, one should re-train and select a new set of discriminative features using profile face images.
 - A key idea of deep learning for visual recognition is to learn a feature hierarchy all the way from pixels to classifier.

2 Backpropagation [40 points]

Consider the following function:

$$f(w, x, y, z) = (w + xy)z \quad \text{where } w = 1, x = -1, y = 2, z = 1$$

Draw the computational graph for this function, and write all the intermediate variable output values as well as their partial derivatives (i.e., $\partial f / \partial (\text{variable})$). Clearly define any intermediate variables that you use (*Hint*: there should be two). Ultimately, we are interested in computing the following four quantities:

$$\partial f / \partial w, \partial f / \partial x, \partial f / \partial y, \partial f / \partial z$$